

```

import pandas as pd
import numpy as np
import seaborn as sns
from sklearn.cluster import KMeans
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report, confusion_matrix

df = pd.read_csv(r"C:\Users\divaa\OneDrive\Desktop\pri\Blend\Blend
dataset\customer_segmentation.csv")
df

```

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	Male	19	15	39
1	2	Male	21	15	81
2	3	Female	20	16	6
3	4	Female	23	16	77
4	5	Female	31	17	40
..
..					
195	196	Female	35	120	
79					
196	197	Female	45	126	
28					
197	198	Male	32	126	
74					
198	199	Male	32	137	
18					
199	200	Male	30	137	
83					

[200 rows x 5 columns]

```
print(df.head())
```

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	Male	19	15	39
1	2	Male	21	15	81
2	3	Female	20	16	6
3	4	Female	23	16	77
4	5	Female	31	17	40

```
print(df.isnull().sum())
```

```
CustomerID      0
Gender           0
Age             0
Annual Income (k$)  0
Spending Score (1-100)  0
dtype: int64
```

```
print(df.info())
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 5 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   CustomerID                            200 non-null    int64
1   Gender                                200 non-null    object
2   Age                                    200 non-null    int64
3   Annual Income (k$)                    200 non-null    int64
4   Spending Score (1-100)                200 non-null    int64
dtypes: int64(4), object(1)
memory usage: 7.9+ KB
None
```

```
df.describe()
```


	CustomerID	Age	Annual Income (k\$)	Spending Score (1-100)
count	200.000000	200.000000	200.000000	200.000000
mean	100.500000	38.850000	60.560000	50.200000
std	57.879185	13.969007	26.264721	25.823522
min	1.000000	18.000000	15.000000	1.000000
25%	50.750000	28.750000	41.500000	34.750000
50%	100.500000	36.000000	61.500000	50.000000
75%	150.250000	49.000000	78.000000	73.000000
max	200.000000	70.000000	137.000000	99.000000

```
plt.figure(figsize=(10, 6))
```


```
sns.histplot(df['Age'], kde=True, bins=30)
plt.title('Age Distribution')
plt.xlabel('Age')
plt.ylabel('Frequency')
plt.show()
```



```
plt.figure(figsize=(10, 6))
sns.histplot(df['Annual Income (k$)'], kde=True, bins=30)
plt.title('Annual Income Distribution')
plt.xlabel('Annual Income (k$)')
plt.ylabel('Frequency')
plt.show()
```



```
plt.figure(figsize=(10, 6))
sns.histplot(df['Spending Score (1-100)'], kde=True, bins=30)
plt.title('Spending Score Distribution')
plt.xlabel('Spending Score (1-100)')
plt.ylabel('Frequency')
plt.show()
```



```

from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
import matplotlib.pyplot as plt
scaler = StandardScaler()
df_scaled = scaler.fit_transform(df[['Annual Income (k$)', 'Spending
Score (1-100)']])
wcss = []
for i in range(1, 11):
    kmeans = KMeans(n_clusters=i, init='k-means++', max_iter=300,
n_init=10, random_state=42, algorithm='elkan')
    kmeans.fit(df_scaled)
    wcss.append(kmeans.inertia_)
plt.figure(figsize=(10, 6))
plt.plot(range(1, 11), wcss, marker='o')
plt.title('Elbow Method')
plt.xlabel('Number of clusters')
plt.ylabel('WCSS')
plt.show()


```

```

C:\Users\divaa\anaconda_new\Lib\site-packages\sklearn\cluster\
_kmeans.py:1418: RuntimeWarning: algorithm='elkan' doesn't make sense
for a single cluster. Using 'lloyd' instead.
  warnings.warn(
C:\Users\divaa\anaconda_new\Lib\site-packages\sklearn\cluster\
_kmeans.py:1429: UserWarning: KMeans is known to have a memory leak on

```

Windows with MKL, when there are less chunks than available threads.
You can avoid it by setting the environment variable
OMP_NUM_THREADS=1.
warnings.warn(



```
df['Gender'] = df['Gender'].map({'Male': 1, 'Female': 0})
X = df[['Age', 'Annual Income (k$)', 'Spending Score (1-100)']]
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_scaled, df['Spending Score (1-100)'])

KNeighborsClassifier()

y_pred = knn.predict(X_scaled)
print(confusion_matrix(df['Spending Score (1-100)'], y_pred))
print(classification_report(df['Spending Score (1-100)'], y_pred))

[[2 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 [0 0 2 ... 0 0 0]
 ...]
```

```
[0 0 0 ... 0 0 0]
[0 0 0 ... 0 0 0]
[0 0 0 ... 0 0 0]]
```

	precision	recall	f1-score	support
1	0.67	1.00	0.80	2
3	0.00	0.00	0.00	1
4	0.67	1.00	0.80	2
5	0.27	0.75	0.40	4
6	1.00	0.50	0.67	2
7	0.33	1.00	0.50	1
8	0.33	1.00	0.50	1
9	0.00	0.00	0.00	1
10	0.00	0.00	0.00	2
11	0.00	0.00	0.00	1
12	0.00	0.00	0.00	1
13	0.17	0.33	0.22	3
14	0.29	0.50	0.36	4
15	0.00	0.00	0.00	3
16	0.00	0.00	0.00	2
17	0.00	0.00	0.00	3
18	0.00	0.00	0.00	1
20	0.50	1.00	0.67	2
22	0.00	0.00	0.00	1
23	0.00	0.00	0.00	1
24	0.00	0.00	0.00	1
26	0.00	0.00	0.00	2
27	0.00	0.00	0.00	1
28	0.00	0.00	0.00	2
29	0.33	0.50	0.40	2
31	0.00	0.00	0.00	1
32	0.00	0.00	0.00	2
34	0.00	0.00	0.00	1
35	0.38	0.60	0.46	5
36	0.00	0.00	0.00	2
39	0.00	0.00	0.00	2
40	0.40	0.50	0.44	4
41	0.50	0.50	0.50	4
42	0.40	0.75	0.52	8
43	0.50	0.67	0.57	3
44	0.00	0.00	0.00	1
45	0.00	0.00	0.00	1
46	0.31	0.67	0.42	6
47	0.00	0.00	0.00	4
48	0.00	0.00	0.00	5
49	0.00	0.00	0.00	3
50	0.25	0.20	0.22	5
51	0.00	0.00	0.00	3
52	0.50	0.40	0.44	5


53	0.00	0.00	0.00	1
54	0.00	0.00	0.00	3
55	0.00	0.00	0.00	7
56	0.00	0.00	0.00	4
57	0.00	0.00	0.00	2
58	0.00	0.00	0.00	2
59	0.00	0.00	0.00	5
60	0.00	0.00	0.00	3
61	0.67	1.00	0.80	2
63	0.00	0.00	0.00	1
65	0.00	0.00	0.00	1
66	0.00	0.00	0.00	1
68	0.00	0.00	0.00	1
69	0.25	0.50	0.33	2
71	0.00	0.00	0.00	1
72	0.00	0.00	0.00	2
73	0.42	0.83	0.56	6
74	0.00	0.00	0.00	2
75	0.50	0.40	0.44	5
76	0.00	0.00	0.00	2
77	0.33	0.67	0.44	3
78	0.00	0.00	0.00	1
79	0.00	0.00	0.00	2
81	0.00	0.00	0.00	2
82	0.00	0.00	0.00	1
83	0.00	0.00	0.00	2
85	0.50	1.00	0.67	1
86	0.50	1.00	0.67	2
87	0.00	0.00	0.00	2
88	0.33	0.33	0.33	3
89	0.00	0.00	0.00	1
90	0.00	0.00	0.00	2
91	0.00	0.00	0.00	2
92	0.67	0.67	0.67	3
93	0.50	1.00	0.67	2
94	0.00	0.00	0.00	1
95	0.00	0.00	0.00	2
97	0.00	0.00	0.00	2
98	0.00	0.00	0.00	1
99	0.00	0.00	0.00	1
accuracy			0.29	200
macro avg	0.15	0.23	0.17	200
weighted avg	0.19	0.29	0.22	200

C:\Users\divaa\anaconda_new\Lib\site-packages\sklearn\metrics_classification.py:1531: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.


```
_warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))
C:\Users\divaa\anaconda_new\Lib\site-packages\sklearn\metrics\
_classification.py:1531: UndefinedMetricWarning: Precision is ill-
defined and being set to 0.0 in labels with no predicted samples. Use
`zero_division` parameter to control this behavior.
_warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))
C:\Users\divaa\anaconda_new\Lib\site-packages\sklearn\metrics\
_classification.py:1531: UndefinedMetricWarning: Precision is ill-
defined and being set to 0.0 in labels with no predicted samples. Use
`zero_division` parameter to control this behavior.
_warn_prf(average, modifier, f"{metric.capitalize()} is",
len(result))
```

```
from sklearn.cluster import KMeans
kmeans = KMeans(n_clusters=5)
clusters = kmeans.fit_predict(X_scaled)
df['Cluster'] = clusters
plt.scatter(X['Annual Income (k$)'], X['Spending Score (1-100)'],
c=df['Cluster'], cmap='rainbow')
plt.xlabel('Annual Income (k$)')
plt.ylabel('Spending Score (1-100)')
plt.show()
```


```
C:\Users\divaa\anaconda_new\Lib\site-packages\sklearn\cluster\
_kmeans.py:1429: UserWarning: KMeans is known to have a memory leak on
Windows with MKL, when there are less chunks than available threads.
You can avoid it by setting the environment variable
OMP_NUM_THREADS=1.
warnings.warn(
```



```
import seaborn as sns
import matplotlib.pyplot as plt
df_numeric = df.drop(columns=['CustomerID', 'Gender']) # Drop non-numeric columns
unique_colors = ['#FF6347', '#1E90FF', '#32CD32', '#FFD700', '#8A2BE2'] # Added an extra color
cluster_summary = df_numeric.groupby('Cluster').mean()[['Age', 'Annual Income (k$)', 'Spending Score (1-100)']]
print(cluster_summary)
plt.figure(figsize=(10, 6))
sns.countplot(x='Cluster', data=df, palette=unique_colors, hue='Cluster', dodge=False)
plt.title('Cluster Sizes')
plt.xlabel('Cluster')
plt.ylabel('Number of Customers')
plt.legend(title='Cluster')
plt.show()
```

	Age	Annual Income (k\$)	Spending Score (1-100)
Cluster			
0	55.275862	47.620690	41.706897
1	44.387097	89.774194	18.483871
2	26.733333	54.311111	40.911111

3	25.769231	26.115385	74.846154
4	32.875000	86.100000	81.525000



```
print(df.head())
```

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	NaN	19	15	39
1	2	NaN	21	15	81
2	3	NaN	20	16	6
3	4	NaN	23	16	77
4	5	NaN	31	17	40

	Cluster
0	3
1	3
2	2
3	3
4	3

```

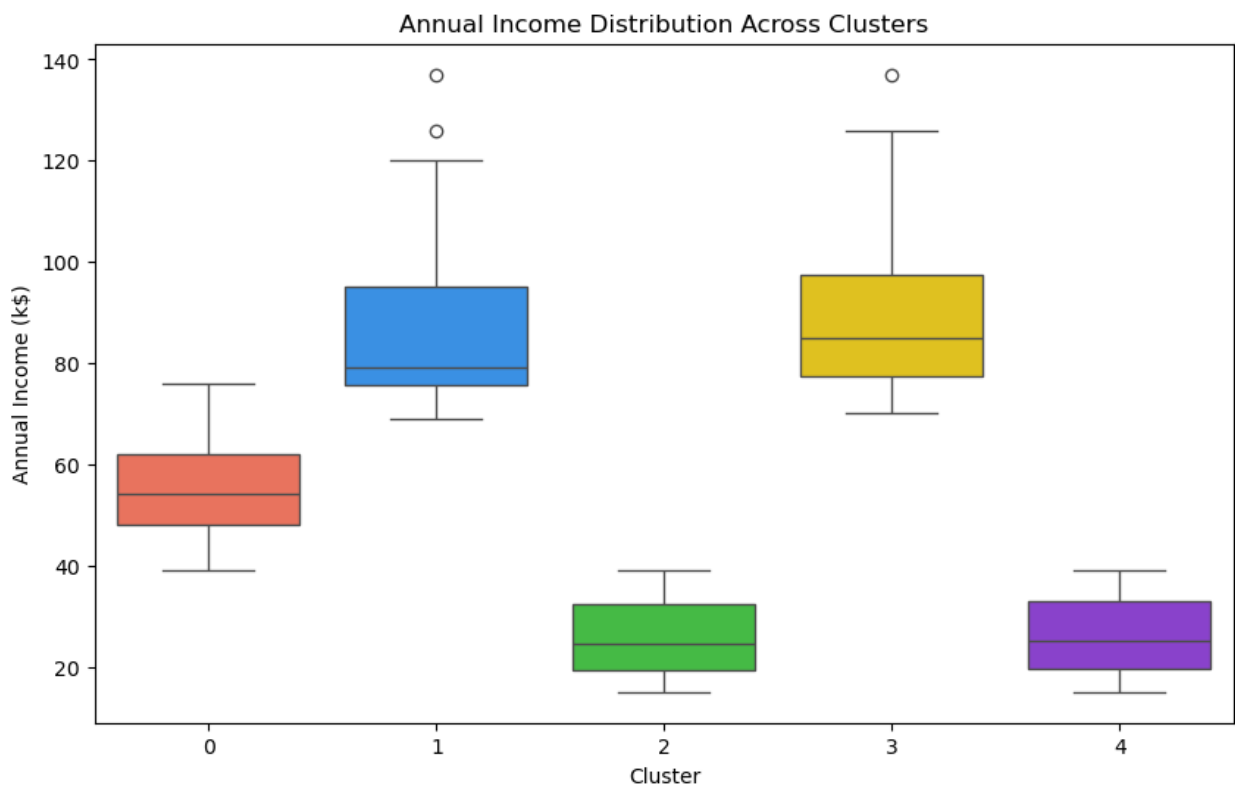
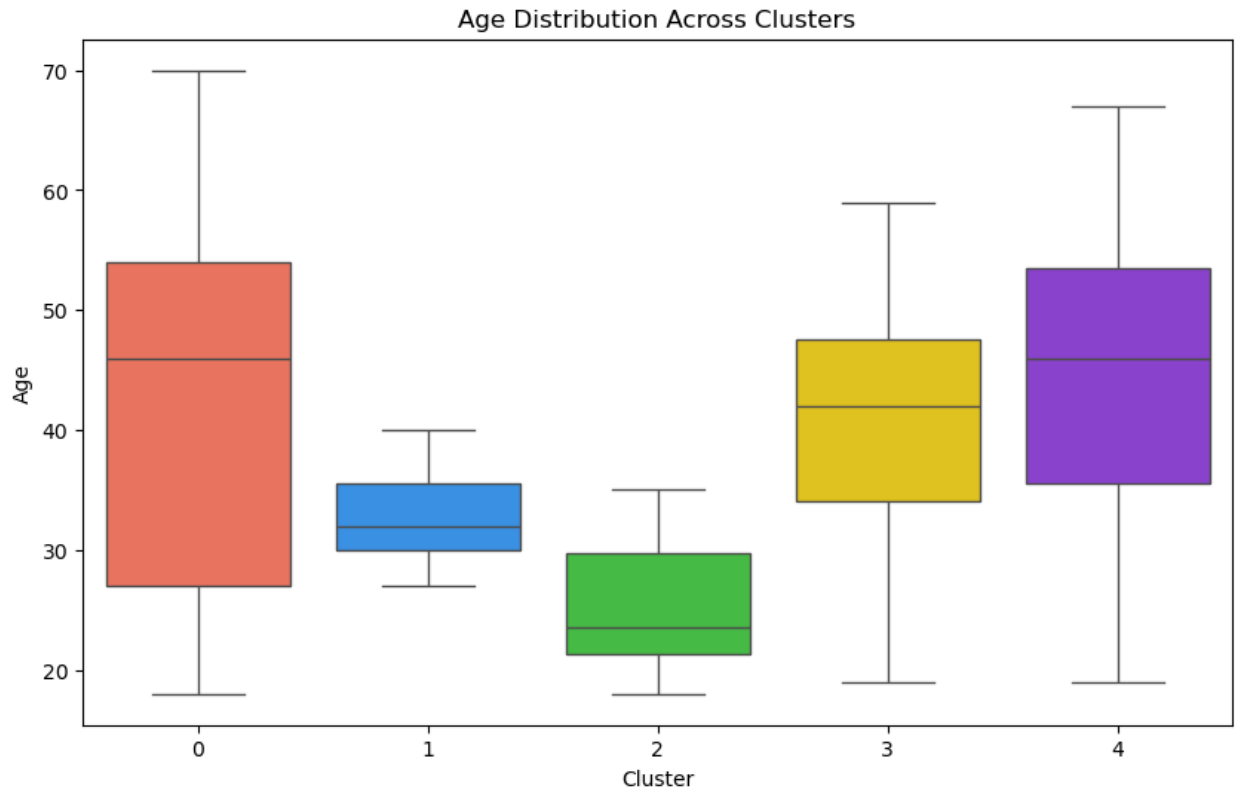
unique_colors = ['#FF6347', '#1E90FF', '#32CD32', '#FFD700',
'#8A2BE2']


# Boxplot for Age distribution across clusters
plt.figure(figsize=(10, 6))
sns.boxplot(x='Cluster', y='Age', data=df, hue='Cluster',
palette=unique_colors, dodge=False)
plt.legend([],[], frameon=False) # Hide legend if not needed
plt.title('Age Distribution Across Clusters')
plt.show()

# Boxplot for Annual Income distribution across clusters
plt.figure(figsize=(10, 6))
sns.boxplot(x='Cluster', y='Annual Income (k$)', data=df,
hue='Cluster', palette=unique_colors, dodge=False)
plt.legend([],[], frameon=False) # Hide legend if not needed
plt.title('Annual Income Distribution Across Clusters')
plt.show()

# Boxplot for Spending Score distribution across clusters
plt.figure(figsize=(10, 6))
sns.boxplot(x='Cluster', y='Spending Score (1-100)', data=df,
hue='Cluster', palette=unique_colors, dodge=False)
plt.legend([],[], frameon=False) # Hide legend if not needed
plt.title('Spending Score Distribution Across Clusters')
plt.show()

```





```
from sklearn.cluster import MiniBatchKMeans
from sklearn.preprocessing import StandardScaler
import seaborn as sns
import matplotlib.pyplot as plt
df_cluster = df[['Age', 'Spending Score (1-100)']] # Or you can use
['Age', 'Annual Income (k$)']
scaler = StandardScaler()
df_scaled = scaler.fit_transform(df_cluster)
minibatch_kmeans = MiniBatchKMeans(n_clusters=5, init='k-means++',
max_iter=300, batch_size=100, random_state=42)
df['Cluster'] = minibatch_kmeans.fit_predict(df_scaled)
unique_colors = ['#FF5733', '#33FF57', '#3357FF', '#FF33A1',
'#A133FF'] # Customize with more unique colors
```

```
plt.figure(figsize=(10, 6))
sns.scatterplot(x='Age', y='Spending Score (1-100)', hue='Cluster',
data=df, palette=unique_colors)
plt.title('Age vs. Spending Score by Cluster')
plt.xlabel('Age')
plt.ylabel('Spending Score (1-100)')
plt.show()
```

```
plt.figure(figsize=(10, 6))
sns.scatterplot(x='Age', y='Annual Income (k$)', hue='Cluster',
```

```

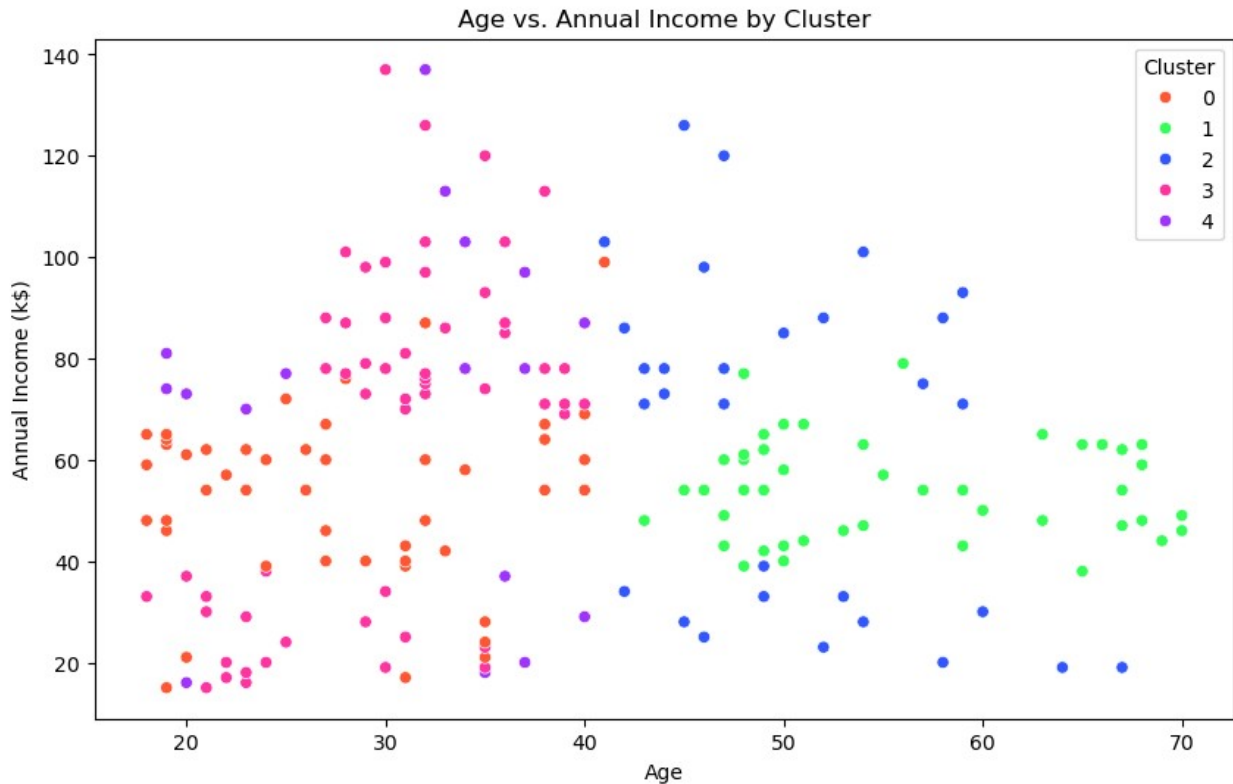
data=df, palette=unique_colors)
plt.title('Age vs. Annual Income by Cluster')
plt.xlabel('Age')
plt.ylabel('Annual Income (k$)')
plt.show()

# Optional: If you want to display the cluster centers
centers = scaler.inverse_transform(minibatch_kmeans.cluster_centers_)
print("Cluster Centers (Age, Spending Score):")
print(centers)

C:\Users\divaa\anaconda_new\Lib\site-packages\sklearn\cluster\
_kmeans.py:1955: UserWarning: MiniBatchKMeans is known to have a
memory leak on Windows with MKL, when there are less chunks than
available threads. You can prevent it by setting batch_size >= 2048 or
by setting the environment variable OMP_NUM_THREADS=1
  warnings.warn(

```





Cluster Centers (Age, Spending Score):

```
[[27.81308411 49.08010681]
 [55.35519802 48.26485149]
 [50.36417323 16.41141732]
 [30.50544016 82.57171118]
 [31.2654321 15.52469136]]
```

Drop non-numeric columns for numerical analysis

```
df_numeric = df.drop(columns=['CustomerID', 'Gender']) # Drop non-numeric columns like 'Gender'
```

Group by cluster and calculate the mean for each cluster

```
cluster_summary = df_numeric.groupby('Cluster').mean()
```

```
print("Cluster Summary (Mean Values):")
```

```
print(cluster_summary)
```

Get the size of each cluster (number of customers in each cluster)

```
cluster_sizes = df_numeric['Cluster'].value_counts().sort_index()
```

```
print("\nCluster Sizes (Number of Customers per Cluster):")
```

```
print(cluster_sizes)
```

If you also want median values for more robustness:

```
cluster_median = df_numeric.groupby('Cluster').median()
```

```
print("\nCluster Summary (Median Values):")
```

```
print(cluster_median)
```


If you want to see additional statistics like standard deviation, min, max, etc., you can use:

```
cluster_stats = df_numeric.groupby('Cluster').describe()
```

```
print("\nCluster Detailed Statistics:")
```

```
print(cluster_stats)
```

Cluster Summary (Mean Values):

	Age	Annual Income (k\$)	Spending Score (1-100)
Cluster			
0	55.275862	47.620690	41.706897
1	44.387097	89.774194	18.483871
2	26.733333	54.311111	40.911111
3	25.769231	26.115385	74.846154
4	32.875000	86.100000	81.525000

Cluster Sizes (Number of Customers per Cluster):

Cluster	
0	58
1	31
2	45
3	26
4	40

Name: count, dtype: int64

Cluster Summary (Median Values):

	Age	Annual Income (k\$)	Spending Score (1-100)
Cluster			
0	53.0	48.5	46.0
1	44.0	87.0	17.0
2	26.0	59.0	46.0
3	24.0	24.5	75.5
4	32.0	78.5	83.0

Cluster Detailed Statistics:

	Age							
	count	mean	std	min	25%	50%	75%	max
Cluster								
0	58.0	55.275862	8.571256	40.0	49.00	53.0	63.75	70.0
1	31.0	44.387097	8.232770	32.0	37.00	44.0	49.00	59.0
2	45.0	26.733333	7.085196	18.0	20.00	26.0	32.00	40.0
3	26.0	25.769231	5.435496	18.0	21.25	24.0	30.75	35.0
4	40.0	32.875000	3.857643	27.0	30.00	32.0	36.00	40.0

	Annual Income (k\$)				
	count	mean	...	75%	max
Cluster			...		
0	58.0	47.620690	...	59.75	67.0
1	31.0	89.774194	...	98.50	137.0
2	45.0	54.311111	...	64.00	81.0
3	26.0	26.115385	...	33.00	42.0

4	40.0	86.100000	...	94.00	137.0
Spending Score (1-100)					
\	count	mean	std	min	25%
50%					
Cluster					
0	58.0	41.706897	15.697814	3.0	37.25
46.0					
1	31.0	18.483871	10.194348	1.0	12.00
17.0					
2	45.0	40.911111	16.285552	5.0	35.00
46.0					
3	26.0	74.846154	15.069684	39.0	67.50
75.5					
4	40.0	81.525000	9.999968	58.0	74.00
83.0					
	75%	max			
Cluster					
0	52.00	60.0			
1	25.00	39.0			
2	54.00	60.0			
3	81.75	99.0			
4	90.00	97.0			
[5 rows x 24 columns]					

Cluster 0: Customers are older (mean age 43), have a moderate income (mean \$55k), and an average spending score (mean 50). This group represents a balanced customer segment in terms of spending and income.

Cluster 1: Customers are younger (mean age 33), have a high income (mean \$87k), and a very high spending score (mean 82). These are affluent and high-spending customers.

Cluster 2: Customers are the youngest (mean age 25), have a low income (mean \$26k), but a high spending score (mean 79). This segment spends more despite lower incomes.

Cluster 3: Customers are middle-aged (mean age 41), have a high income (mean \$88k), but a very low spending score (mean 17). These high-income customers are conservative spenders.

Cluster 4: Customers are older (mean age 45), have a low income (mean \$26k), and a low spending score (mean 21). This group has both low income and low spending behavior.

The customer segmentation reveals five distinct groups with varying age, income, and spending behaviors. Affluent, high-spending customers (Cluster 1) should be targeted with premium offers, while younger, budget-conscious groups (Clusters 2 and 4) respond better to affordable, trendy products. Middle-aged, high-income but conservative spenders (Cluster 3) need value-focused strategies to increase engagement.

